

Importing Data/Packages

```
In [2]: import pandas as pd
import numpy as np
from sklearn.tree import DecisionTreeClassifier # Import Decision Tree Classifier
from sklearn.model_selection import train_test_split # Import train_test_split function
from sklearn import metrics #Import scikit-learn metrics module for accuracy calculation
```

```
In [3]: dataset= pd.read_csv('https://raw.githubusercontent.com/pleunipennings/EDSS/df328799ceb8dc834c63669645f79ee44dbflabc/allfeaturesV8_EDSS.csv')
dataset
```

Out[3]:

	pos	makesCpG	bigAAChange	Avg_Mutation_Freq	High_Mutation_Freq	5' UTR	Core	E1
0	8621	0	0	0.000009	0	0	0	0
1	8622	0	1	0.000019	0	0	0	0
2	8627	0	0	0.000019	0	0	0	0
3	8633	0	0	0.000024	0	0	0	0
4	8634	1	1	0.000024	0	0	0	0
...
7879	3191	0	0	0.299070	1	0	0	0
7880	3473	0	0	0.299204	1	0	0	0
7881	1512	0	1	0.299441	1	0	0	0
7882	6971	0	0	0.299442	1	0	0	0
7883	800	1	0	0.299490	1	0	1	0

7884 rows × 29 columns

In [4]: `dataset.describe()` *#provides dataframe overview*

Out[4]:

	pos	makesCpG	bigAACChange	Avg_Mutation_Freq	High_Mutation_Freq	5' UTR
count	7884.000000	7884.000000	7884.000000	7884.000000	7884.000000	7884.000000
mean	4452.503171	0.146372	0.503425	0.031108	0.504693	0.000000
std	2427.644331	0.353502	0.500020	0.055110	0.500010	0.090000
min	264.000000	0.000000	0.000000	0.000009	0.000000	0.000000
25%	2346.750000	0.000000	0.000000	0.002362	0.000000	0.000000
50%	4460.500000	0.000000	1.000000	0.004376	1.000000	0.000000
75%	6552.250000	0.000000	1.000000	0.033633	1.000000	0.000000
max	8641.000000	1.000000	1.000000	0.299490	1.000000	1.000000

8 rows × 29 columns

In [5]: `dataset.dtypes` *#can see data classes for each parameter in data*

Out[5]:

pos	int64
makesCpG	int64
bigAACChange	int64
Avg_Mutation_Freq	float64
High_Mutation_Freq	int64
5' UTR	int64
Core	int64
E1	int64
E2	int64
HVR1	int64
NS1	int64
NS2	int64
NS3	int64
NS4A	int64
NS4B	int64
NS5A	int64
NS5B	int64
Mutation_Rate	float64
RNAstructure	int64
Nonsyn	int64
Positive AA	int64
Negative AA	int64
Hydrophobic AA	int64
Polar AA	int64
Nonpolar AA	int64
Acidic AA	int64
Basic AA	int64
Stop	int64
Drastic	int64
dtype:	object

Looking for any missing data

```
In [6]: (dataset=='?').sum() #no missing data found
```

```
Out[6]: pos                0
        makesCpG           0
        bigAACChange       0
        Avg_Mutation_Freq  0
        High_Mutation_Freq 0
        5' UTR             0
        Core               0
        E1                 0
        E2                 0
        HVR1              0
        NS1                0
        NS2                0
        NS3                0
        NS4A              0
        NS4B              0
        NS5A              0
        NS5B              0
        Mutation_Rate      0
        RNAstructure       0
        Nonsyn             0
        Positive AA        0
        Negative AA        0
        Hydrophobic AA     0
        Polar AA           0
        Nonpolar AA        0
        Acidic AA          0
        Basic AA           0
        Stop               0
        Drastic            0
        dtype: int64
```

Split target and features

```
In [7]: features = dataset.drop(columns='High_Mutation_Freq') #isolate column of
data from dataframe to use for prediction/analysis
features2 = features.drop(columns="Avg_Mutation_Freq")
features2
```

Out[7]:

	pos	makesCpG	bigAACChange	5' UTR	Core	E1	E2	HVR1	NS1	NS2	...	Nonsyn	Posit
0	8621	0	0	0	0	0	0	0	0	0	...	0	
1	8622	0	1	0	0	0	0	0	0	0	...	1	
2	8627	0	0	0	0	0	0	0	0	0	...	0	
3	8633	0	0	0	0	0	0	0	0	0	...	0	
4	8634	1	1	0	0	0	0	0	0	0	...	1	
...	
7879	3191	0	0	0	0	0	0	0	0	1	...	0	
7880	3473	0	0	0	0	0	0	0	0	0	...	0	
7881	1512	0	1	0	0	0	0	1	0	0	...	1	
7882	6971	0	0	0	0	0	0	0	0	0	...	0	
7883	800	1	0	0	1	0	0	0	0	0	...	0	

7884 rows × 27 columns

```
In [8]: labels = dataset["High_Mutation_Freq"] #saves target columns into variable, "labels"
print(labels)
```

```
0      0
1      0
2      0
3      0
4      0
..
7879   1
7880   1
7881   1
7882   1
7883   1
```

Name: High_Mutation_Freq, Length: 7884, dtype: int64

```
In [9]: labels = np.where(labels >= 1,1,0) #convert labels to binary values that
show either high or low mutation frequency
labels
```

Out[9]: array([0, 0, 0, ..., 1, 1, 1])

```
In [10]: print("Number of sites with low mutation frequency = " + str(np.count_nonzero(labels==0))) #organizing data into number of sites with high or low mutation frequencies
print("Number of sites with high mutation frequency = " + str(np.count_nonzero(labels==1)))
```

Number of sites with low mutation frequency = 3905
 Number of sites with high mutation frequency = 3979

Train vs. Test Data

```
In [11]: features2_train, features2_test, labels_train, labels_test = train_test_split(features2, labels, test_size=0.3, random_state=42) #random_state=  
randomizing the data in a standard way for everyone  
# 70% training and 30% test
```

```
In [12]: # Create Decision Tree classifier object  
clf = DecisionTreeClassifier(max_depth = 4)  
#max_depth: no more than x number of questions  
  
# Train Decision Tree Classifier  
clf = clf.fit(features2_train, labels_train )  
# .fit function creates the decision tree
```

Predictions

```
In [13]: #Predict the response for test dataset  
labels_pred = clf.predict(features2_test)  
labels_pred
```

```
Out[13]: array([0, 1, 0, ..., 0, 1, 1])
```

```
In [14]: # Look at the predicted values. Remember, 0 means no constricted vessels, 1 means at least one.  
print(labels_pred)  
# And the real values.  
print(labels_test)
```

```
[0 1 0 ... 0 1 1]  
[0 1 1 ... 0 1 0]
```


Results/Confusion Matrix

```
In [15]: print(metrics.confusion_matrix(labels_test, labels_pred)) #There is 100%
in prediction of sites of high and low frequency mutations.
metrics.plot_confusion_matrix(clf, features2_test, labels_test)
```

```
[[960 210]
 [272 924]]
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87:
FutureWarning: Function plot_confusion_matrix is deprecated; Function `
plot_confusion_matrix` is deprecated in 1.0 and will be removed in 1.2.
Use one of the class methods: ConfusionMatrixDisplay.from_predictions o
r ConfusionMatrixDisplay.from_estimator.
warnings.warn(msg, category=FutureWarning)
```

```
Out[15]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f2
444c74fd0>
```



```
In [16]: acc = round(100 * metrics.accuracy_score(labels_test, labels_pred),2)
print("Accuracy:", acc, "%")
```

```
Accuracy: 79.63 %
```

Decision Tree

```
In [17]: from matplotlib import pyplot as plt
from sklearn import tree
```

```
In [18]: features2.columns[:-1]
```

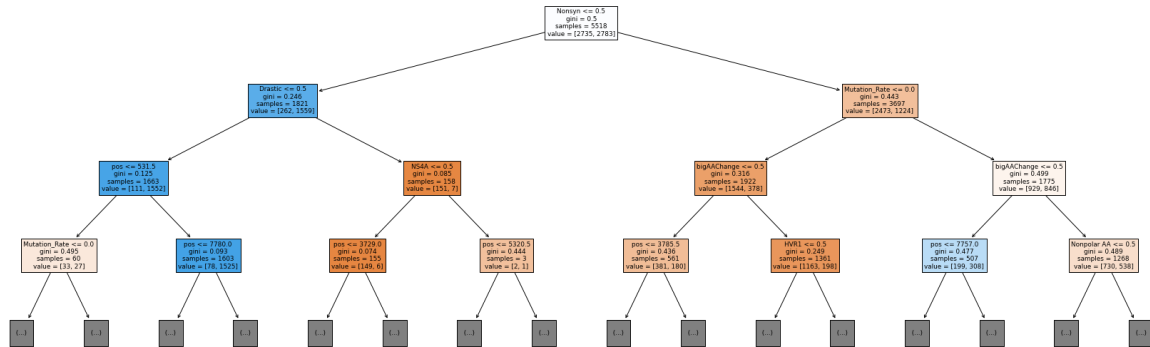
```
Out[18]: Index(['pos', 'makesCpG', 'bigAAChange', '5' UTR', 'Core', 'E1', 'E2',
'HVR1',
'NS1', 'NS2', 'NS3', 'NS4A', 'NS4B', 'NS5A', 'NS5B', 'Mutation_R
ate',
'RNAstructure', 'Nonsyn', 'Positive AA', 'Negative AA',
'Hydrophobic AA', 'Polar AA', 'Nonpolar AA', 'Acidic AA', 'Basic
AA',
'Stop'],
dtype='object')
```

```
In [19]: fig = plt.figure(figsize=(30,10)) #decision tree  
         tree.plot_tree(clf, filled=True, max_depth=3, feature_names = features2.  
                        columns)
```

```

Out[19]: [Text(0.5, 0.9, 'Nonsyn <= 0.5\ngini = 0.5\nsamples = 5518\nvalue = [27
35, 2783]'),
Text(0.25, 0.7, 'Drastic <= 0.5\ngini = 0.246\nsamples = 1821\nvalue =
[262, 1559]'),
Text(0.125, 0.5, 'pos <= 531.5\ngini = 0.125\nsamples = 1663\nvalue =
[111, 1552]'),
Text(0.0625, 0.3, 'Mutation_Rate <= 0.0\ngini = 0.495\nsamples = 60\nv
alue = [33, 27]'),
Text(0.03125, 0.1, '\n (...) \n'),
Text(0.09375, 0.1, '\n (...) \n'),
Text(0.1875, 0.3, 'pos <= 7780.0\ngini = 0.093\nsamples = 1603\nvalue
= [78, 1525]'),
Text(0.15625, 0.1, '\n (...) \n'),
Text(0.21875, 0.1, '\n (...) \n'),
Text(0.375, 0.5, 'NS4A <= 0.5\ngini = 0.085\nsamples = 158\nvalue = [1
51, 7]'),
Text(0.3125, 0.3, 'pos <= 3729.0\ngini = 0.074\nsamples = 155\nvalue =
[149, 6]'),
Text(0.28125, 0.1, '\n (...) \n'),
Text(0.34375, 0.1, '\n (...) \n'),
Text(0.4375, 0.3, 'pos <= 5320.5\ngini = 0.444\nsamples = 3\nvalue =
[2, 1]'),
Text(0.40625, 0.1, '\n (...) \n'),
Text(0.46875, 0.1, '\n (...) \n'),
Text(0.75, 0.7, 'Mutation_Rate <= 0.0\ngini = 0.443\nsamples = 3697\nv
alue = [2473, 1224]'),
Text(0.625, 0.5, 'bigAAChange <= 0.5\ngini = 0.316\nsamples = 1922\nva
lue = [1544, 378]'),
Text(0.5625, 0.3, 'pos <= 3785.5\ngini = 0.436\nsamples = 561\nvalue =
[381, 180]'),
Text(0.53125, 0.1, '\n (...) \n'),
Text(0.59375, 0.1, '\n (...) \n'),
Text(0.6875, 0.3, 'HVR1 <= 0.5\ngini = 0.249\nsamples = 1361\nvalue =
[1163, 198]'),
Text(0.65625, 0.1, '\n (...) \n'),
Text(0.71875, 0.1, '\n (...) \n'),
Text(0.875, 0.5, 'bigAAChange <= 0.5\ngini = 0.499\nsamples = 1775\nva
lue = [929, 846]'),
Text(0.8125, 0.3, 'pos <= 7757.0\ngini = 0.477\nsamples = 507\nvalue =
[199, 308]'),
Text(0.78125, 0.1, '\n (...) \n'),
Text(0.84375, 0.1, '\n (...) \n'),
Text(0.9375, 0.3, 'Nonpolar AA <= 0.5\ngini = 0.489\nsamples = 1268\nv
alue = [730, 538]'),
Text(0.90625, 0.1, '\n (...) \n'),
Text(0.96875, 0.1, '\n (...) \n')]

```

Lab Reflection

Questions

Describe what you did (describe the steps you took in the notebook) What you found (describe your conclusions from the decision tree; which features are most important in predicting the mutation frequency? what is the accuracy of your prediction?).

Write 300-400 words (as a team).

We took various steps, such as reading the Hep C data csv file into this notebook, checking for any gaps of missing data in the dataset, removing the target column of data from the dataset, etc. We also made sure to remove average mutation frequency, as this could hinder what/how predictions were made. We converted the column into binary data, and split it into 70% train and 30% test data. Then, we created a decision tree and confusion matrix. The confusion matrix and accuracy functions showed the model was 79.63% accurate at correctly predicting high and low frequency mutations.

The decision tree separated the data into synonymous vs. non-synonymous mutations. (Synonymous mutations result in no change to the amino acid being coded, whereas non-synonymous mutations results in a change to the amino acid and therefore the protein formed.) Then, on the right side, it separated into high vs. low frequency mutations. Low frequency mutations are more of a threat to the virus's survival, as opposed to high frequency mutations. On the right side, the it was further divided into how big of a change the mutations resulted in to affect the amino acid being coded.

The higher the certainty is about the predictions, the bolder the colors are. This can be seen towards the bottom of the decision tree. The gini index also gets lower towards the bottom of the tree to represent more unbalanced data. It also appears that the synonymous mutations were more accurately predicted as opposed to non-synonymous mutations, based on how bold the colors are in the decision tree.